

Deep Learning

CISC 867

Project 1

By:

Name : Abdelrahman Magdy El-Hamoly

email: 21amme1@queensu.ca

Name : Kareem Gamal Mahmoud Mohamed

email: 21kgmm@queensu.ca

Name : Ahmed Mohamed Gaber

email: 21amga@queensu.ca


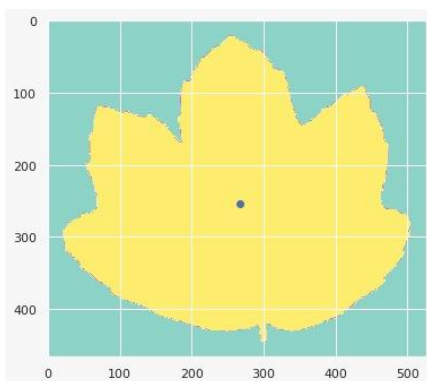
Supervised by\ Dr. Hazem Abbas

Part 1 :


Description and cleaning:

The data is clear ,there are no missing values , there are no outliers, the data is normalized because the range of the values between zero and 1, there is no duplication.

Draw some of the images:



The distribution of the target points:



Part 2 : Simple MLP Model

We decided to tune the following hyperparameters:

- Batch Size
- Hidden Nodes Size
- Drop Rate
- Optimizer

In the beginning we had built a function named 'base_line_model' to build several models with different hyperparameters

This function has default hyperparameters which are

[optim = Adam() , bat_size = 32, hid_nodes = 512, drop_rate = 0.5]

We will build our deep learning model into a function to make it easier to use it multiple times

```
#build a base model as function to call this function with different hyperparameter
#but we have default hyperparameter
def base_line_model(optim = Adam() , bat_size = 32, hid_nodes = 512, drop_rate = 0.5):
    # define the keras model
    model = Sequential()
    # In layer_1 our activation function is 'tanh' with default 512 neurons and kernel_initializer 'glorot_uniform'
    model.add(Dense(hid_nodes, activation='tanh', input_shape=(input_features,), kernel_initializer = 'glorot_uniform', bias_initializer='zeros', name = 'Layer_1'))
    #dropout some nerouns to avoid overfitting
    model.add(Dropout(drop_rate))
    #output layer with softmax activation function and has 99 nodes for output shape
    model.add(Dense(99 , activation='softmax', name = 'Output'))
    #compile the model with sparse_categorical_crossentropy loss function and accuracy metrics
    model.compile(optimizer = optim ,loss='sparse_categorical_crossentropy' , metrics=['accuracy'])

    #fit the model with 100 epoch
    history = model.fit(XTrain , yTrain , epochs=100 , batch_size=bat_size , validation_data=(X_val, y_val))

    #return the training model and the history
    return model, history
```

Tune with different optimizers.

Trial 1: Adam optimizer


We had used adam optimizer in our first model with 32 batch size and 0.5 drop out ratio

```
#the source of why we use this hyperparameter for Adam
#https://keras.io/api/optimizers/
lr_schedule = keras.optimizers.schedules.ExponentialDecay(initial_learning_rate=1e-3, decay_steps=10000, decay_rate=0.9)
adam = Adam(learning_rate = lr_schedule)
#call the model function with adam optimizer
model_1, history_1 = base_line_model(adam)
#display summary for the model
model_1.summary()
#evaluate the model
model_1.evaluate(xTest, yTest)
```

Layer (type)	Output Shape	Param #
Layer_1 (Dense)	(None, 512)	98816
dropout (Dropout)	(None, 512)	0
Output (Dense)	(None, 99)	50787

=====
Total params: 149,603
Trainable params: 149,603
Non-trainable params: 0
=====
7/7 [=====] - 0s 3ms/step - loss: 0.2731 - accuracy: 0.9495
[0.27308189868927, 0.9494949579238892]

Loss ratio for
validation & training



Note : as we can see from this graph we can run only 85 epochs instead of 100 epochs

Trial 2: RMS Prop optimizer


Our second model with RMSProp optimizer with 32 batch size and 0.5 drop out ratio.

```
[ ] #the source of why we use this hyperparameter for RMSprop
    #https://keras.io/api/optimizers/rmsprop/
    rms_prop = tf.keras.optimizers.RMSprop(learning_rate=0.01, rho=0.9, momentum=0.0, epsilon=1e-07, centered=False, name="RMSprop")
    #call the model function with RMSprop optimizer
    model_2.history_2 = base_line_model(rms_prop)
    #display summary for the model
    model_2.summary()
    #evaluate the model
    model_2.evaluate(xTest, yTest)
```

Layer (type)	Output Shape	Param #
Layer_1 (Dense)	(None, 512)	98816
dropout_1 (Dropout)	(None, 512)	0
Output (Dense)	(None, 99)	50787

=====
Total params: 149,603
Trainable params: 149,603
Non-trainable params: 0
=====
7/7 [=====] - 0s 4ms/step - loss: 0.2390 - accuracy: 0.9545
[0.23900854587554932, 0.9545454382896423]

Loss ratio for
validation & training



Note : as we can see from this graph we can run only 25 epochs instead of 100 epochs

Trial 3: SGD optimizer

Our second model with SGD optimizer with 32 batch size and 0.5 drop out ratio.


```
#the source of why we use this hyperparameter for RMSprop
#https://keras.io/api/optimizers/sgd/
sgd = tf.keras.optimizers.SGD(learning_rate=0.9, momentum=0.0, nesterov=False, name="SGD")
#
model_3, history_3 = base_line_model(sgd)

model_3.summary()
model_3.evaluate(xTest, yTest)
```

Layer (type)	Output Shape	Param #
Layer_1 (Dense)	(None, 512)	98816
dropout_2 (Dropout)	(None, 512)	0
Output (Dense)	(None, 99)	50787

=====
Total params: 149,603
Trainable params: 149,603
Non-trainable params: 0
=====
7/7 [=====] - 0s 4ms/step - loss: 0.4287 - accuracy: 0.9091
[0.42866960167884827, 0.9090909361839294]

Loss ratio for
validation & training



as we can see the SGD is very bad because the curve is not stable at all

So, from all of that the **RMSprop** is the best one

Now we will use the RMSprop optimizer with different batch size.

Trial 4: with Batch size = 32

So from the previous trails we can see that the RMSprop is the best one

now we will try RMSprop with different batch_size

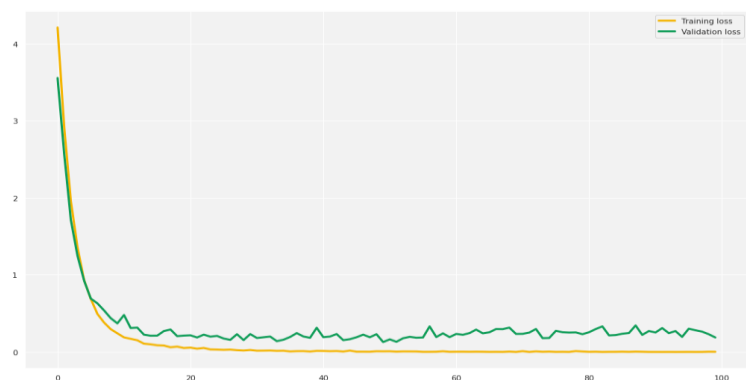
```
#batch_size 32
#rms_prop = tf.keras.optimizers.RMSprop(learning_rate=0.001, rho=0.9, momentum=0.0, epsilon=1e-07, centered=False, name="RMSprop")
#
model_4.history_4 = base_line_model(rms_prop , 32)
#
model_4.summary()
model_4.evaluate(xTest, yTest)
```

Layer (type)	Output Shape	Param #
Layer_1 (Dense)	(None, 512)	98816
dropout_11 (Dropout)	(None, 512)	0
Output (Dense)	(None, 99)	50787

Total params: 149,603
Trainable params: 149,603
Non-trainable params: 0

7/7 [=====] - 0s 5ms/step - loss: 0.2119 - accuracy: 0.9596
[0.21190235018730164, 0.9595959782600403]

Loss ratio for
validation & training



Note : as we can see from this graph we can run only 25 epochs instead of 100 epochs.

Trial 5: with Batch size = 64

now we will try RMSprop with different batch_size


```
[ ] #batch size 64
#rms_prop = tf.keras.optimizers.RMSprop(learning_rate=0.001, rho=0.9, momentum=0.0, epsilon=1e-07, centered=False, name="RMSprop")
#
model_4, history_4 = base_line_model(rms_prop, 64)

model_4.summary()
model_4.evaluate(xTest, yTest)
```

Layer (type)	Output Shape	Param #
Layer_1 (Dense)	(None, 512)	98816
dropout_3 (Dropout)	(None, 512)	0
Output (Dense)	(None, 99)	50787

=====
Total params: 149,603
Trainable params: 149,603
Non-trainable params: 0
=====
7/7 [=====] - 0s 4ms/step - loss: 0.1561 - accuracy: 0.9596
[0.15607938170433044, 0.9595959782600403]

Loss ratio for
validation & training



Note : as we can see from this graph we can run only 25 epochs instead of 100 epochs


Trial 6: with Batch size = 128

```
#batch size 124
#rms_prop = tf.keras.optimizers.RMSprop(learning_rate=0.001,rho=0.9,momentum=0.0,epsilon=1e-07,centered=False, name="RMSprop")
#
model_5,history_5 = base_line_model(rms_prop , 128)
|
model_5.summary()
model_5.evaluate(xTest, yTest)
```

Layer (type)	Output Shape	Param #
Layer_1 (Dense)	(None, 512)	98816
dropout_4 (Dropout)	(None, 512)	0
Output (Dense)	(None, 99)	50787

=====
Total params: 149,603
Trainable params: 149,603
Non-trainable params: 0
=====
7/7 [=====] - 0s 3ms/step - loss: 0.1700 - accuracy: 0.9545
[0.17003951966762543, 0.9545454382896423]

Loss ratio for
validation & training



Note : as we can see from this graph we can run only 25 epochs instead of 100 epochs

So, from all of that the **RMSprop** is the best optimizer and the **batch size = 64** is the best model's hyperparameter Now we will use the RMSprop optimizer and batch size 64 with different number of hidden nodes.

Trial 7: with 1024 hidden nodes

So from the previous trails we can see that the model with RMSprop optimizer and batch_size 64 is the best one


now we will try this model with different numbers of nodes

```
[ ] model_6,history_6 = base_line_model(rms_prop , 64, 1024)

model_6.summary()
model_6.evaluate(xTest, yTest)
```

```
Layer (type)                 Output Shape              Param #
=====
Layer_1 (Dense)              (None, 1024)              197632
dropout_5 (Dropout)          (None, 1024)              0
Output (Dense)               (None, 99)                101475
=====
Total params: 299,107
Trainable params: 299,107
Non-trainable params: 0
=====
7/7 [=====] - 0s 8ms/step - loss: 0.1881 - accuracy: 0.9646
[0.18810120224952698, 0.9646464586257935]
```

Loss ratio for
validation & training



Note : as we can see from this graph we can run only 20 epochs instead of 100 epochs

Trial 8: with 2048 hidden nodes


```
model_7,history_7 = base_line_model(rms_prop , 64, 2048)

model_7.summary()
model_7.evaluate(xTest, yTest)
```

Layer (type)	Output Shape	Param #
Layer_1 (Dense)	(None, 2048)	395264
dropout_6 (Dropout)	(None, 2048)	0
Output (Dense)	(None, 99)	202851

=====
Total params: 598,115
Trainable params: 598,115
Non-trainable params: 0
=====
7/7 [=====] - 0s 4ms/step - loss: 0.2633 - accuracy: 0.9545
[0.263296902179718, 0.9545454382896423]

Loss ratio for
validation & training



Note : as we can see from this graph we can run only 20 epochs instead of 100 epochs

Trial 9: with 512 hidden nodes


```
model_8,history_8 = base_line_model(rms_prop , 64, 512)

model_8.summary()
model_8.evaluate(xTest, yTest)
```

Layer (type)	Output Shape	Param #
Layer_1 (Dense)	(None, 512)	98816
dropout_7 (Dropout)	(None, 512)	0
Output (Dense)	(None, 99)	50787

=====
Total params: 149,603
Trainable params: 149,603
Non-trainable params: 0
=====
7/7 [=====] - 0s 3ms/step - loss: 0.1622 - accuracy: 0.9495
[0.16216205060482025, 0.9494949579238892]

Loss ratio for
validation & training



Note : as we can see from this graph we can run only 25 epochs instead of 100 epochs

So, from all of that the **RMSprop** is the best optimizer and the **batch size = 64** and **the number of hidden nodes are 1024 nodes**, that is the best model hyperparameters. Now we will use the RMSprop optimizer and batch size 64 with 1024 hidden nodes and with different drop out rate .

Trial 10: Drop out rate = 0.3

So from the previous trails we can see that the model with RMSprop optimizer and batch_size 64 and 1024 nodes is the best one

now we will try this model with different numbers of dropout rate


```
model_9.history_9 = base_line_model(rms_prop , 64, 1024, 0.3)

model_9.summary()
model_9.evaluate(xTest, yTest)
```

Layer (type)	Output Shape	Param #
Layer_1 (Dense)	(None, 1024)	197632
dropout_8 (Dropout)	(None, 1024)	0
Output (Dense)	(None, 99)	101475

=====
 Total params: 299,107
 Trainable params: 299,107
 Non-trainable params: 0
 =====
 7/7 [=====] - 0s 5ms/step - loss: 0.2946 - accuracy: 0.9545
 [0.2945918142795563, 0.9545454382896423]

Loss ratio for
validation & training



Note : as we can see from this graph we can run only 20 epochs instead of 100 epochs

Trial 11: Drop out rate = 0.6


```
model_10,history_10 = base_line_model(rms_prop , 64, 1024, 0.6)

model_10.summary()
model_10.evaluate(xTest, yTest)
```

Layer (type)	Output Shape	Param #
Layer_1 (Dense)	(None, 1024)	197632
dropout_9 (Dropout)	(None, 1024)	0
Output (Dense)	(None, 99)	101475

=====
Total params: 299,107
Trainable params: 299,107
Non-trainable params: 0
=====
7/7 [=====] - 0s 3ms/step - loss: 0.1660 - accuracy: 0.9596
[0.16601893305778503, 0.9595959782600403]

Loss ratio for
validation & training



Note : as we can see from this graph we can run only 25 epochs instead of 100 epochs

Trial 12: finally this is the best model with the best hyperparameters :


```
model_11,history_11 = base_line_model(rms_prop , 64, 1024, 0.5)

model_11.summary()
model_11.evaluate(xTest, yTest)
```

- Batch Size = 64
- Hidden Nodes Size = 1024
- Drop Rate = 0.5
- Optimizer = RMSprop

Layer (type)	Output Shape	Param #
Layer_1 (Dense)	(None, 1024)	197632
dropout_10 (Dropout)	(None, 1024)	0
Output (Dense)	(None, 99)	101475

=====
Total params: 299,107
Trainable params: 299,107
Non-trainable params: 0
=====
7/7 [=====] - 0s 4ms/step - loss: 0.1812 - accuracy: 0.9646
[0.18121682107448578, 0.9646464586257935]



so the best model is model_11 =
base_line_model(rms_prop , 64, 1024, 0.5) which has
accuracy 0.9646 on test set and **0.1812 loss** on test set
also and has 100% accuracy on training data

CNN Model

We decided to tune the following hyperparameters:

- Batch Size
- Hidden Nodes Size
- Dropout Rate
- Optimizer

Let's Start with Batch Size. We Used the following values:

1- Batch_size=5

```
model.fit(X_train_r, y_train, epochs=15, validation_data=(X_valid_r, y_valid), batch_size=5)
```

```
Epoch 1/15
159/159 [=====] - 12s 9ms/step - loss: 1.5562 - accuracy: 0.6490 - val_loss: 0.3191 - val_accuracy: 0.8838
Epoch 2/15
159/159 [=====] - 1s 7ms/step - loss: 0.0997 - accuracy: 0.9684 - val_loss: 0.1548 - val_accuracy: 0.9747
Epoch 3/15
159/159 [=====] - 1s 7ms/step - loss: 0.0336 - accuracy: 0.9924 - val_loss: 0.0868 - val_accuracy: 0.9798
Epoch 4/15
159/159 [=====] - 1s 7ms/step - loss: 0.0641 - accuracy: 0.9823 - val_loss: 0.2581 - val_accuracy: 0.9495
Epoch 5/15
159/159 [=====] - 1s 7ms/step - loss: 0.1520 - accuracy: 0.9672 - val_loss: 0.2500 - val_accuracy: 0.9343
Epoch 6/15
159/159 [=====] - 1s 7ms/step - loss: 0.1624 - accuracy: 0.9495 - val_loss: 0.3640 - val_accuracy: 0.9091
Epoch 7/15
159/159 [=====] - 1s 7ms/step - loss: 0.0659 - accuracy: 0.9836 - val_loss: 0.0704 - val_accuracy: 0.9848
Epoch 8/15
159/159 [=====] - 1s 7ms/step - loss: 0.0113 - accuracy: 0.9962 - val_loss: 0.0698 - val_accuracy: 0.9899
Epoch 9/15
159/159 [=====] - 1s 7ms/step - loss: 0.0043 - accuracy: 0.9975 - val_loss: 0.1351 - val_accuracy: 0.9697
Epoch 10/15
159/159 [=====] - 1s 7ms/step - loss: 0.0378 - accuracy: 0.9886 - val_loss: 0.1318 - val_accuracy: 0.9646
Epoch 11/15
159/159 [=====] - 1s 7ms/step - loss: 0.1240 - accuracy: 0.9760 - val_loss: 0.3627 - val_accuracy: 0.9444
Epoch 12/15
159/159 [=====] - 1s 7ms/step - loss: 0.0732 - accuracy: 0.9823 - val_loss: 0.1649 - val_accuracy: 0.9697
Epoch 13/15
159/159 [=====] - 1s 7ms/step - loss: 0.1219 - accuracy: 0.9735 - val_loss: 0.5393 - val_accuracy: 0.9192
Epoch 14/15
159/159 [=====] - 1s 7ms/step - loss: 0.0869 - accuracy: 0.9785 - val_loss: 0.2993 - val_accuracy: 0.9596
Epoch 15/15
159/159 [=====] - 1s 7ms/step - loss: 0.0348 - accuracy: 0.9899 - val_loss: 0.1637 - val_accuracy: 0.9394
<keras.callbacks.History at 0x7f4be96a6d50>
```

```
loss: 0.0348 - accuracy: 0.9899 - val_loss: 0.1637 - val_accuracy: 0.9394
```


2- Batch_size=10

```
model.fit(X_train_r, y_train, epochs=15, validation_data=(X_valid_r, y_valid), batch_size=10)
```

```
Epoch 1/15
80/80 [=====] - 2s 13ms/step - loss: 1.5788 - accuracy: 0.6503 - val_loss: 0.1747 - val_accuracy: 0.9495
Epoch 2/15
80/80 [=====] - 1s 11ms/step - loss: 0.0736 - accuracy: 0.9861 - val_loss: 0.1286 - val_accuracy: 0.9798
Epoch 3/15
80/80 [=====] - 1s 11ms/step - loss: 0.0297 - accuracy: 0.9937 - val_loss: 0.1117 - val_accuracy: 0.9848
Epoch 4/15
80/80 [=====] - 1s 7ms/step - loss: 0.0104 - accuracy: 0.9987 - val_loss: 0.0367 - val_accuracy: 0.9949
Epoch 5/15
80/80 [=====] - 1s 8ms/step - loss: 0.0017 - accuracy: 1.0000 - val_loss: 0.0288 - val_accuracy: 1.0000
Epoch 6/15
80/80 [=====] - 1s 8ms/step - loss: 0.0012 - accuracy: 1.0000 - val_loss: 0.0265 - val_accuracy: 1.0000
Epoch 7/15
80/80 [=====] - 1s 7ms/step - loss: 9.4765e-04 - accuracy: 1.0000 - val_loss: 0.0253 - val_accuracy: 1.0000
Epoch 8/15
80/80 [=====] - 1s 8ms/step - loss: 7.8781e-04 - accuracy: 1.0000 - val_loss: 0.0243 - val_accuracy: 0.9949
Epoch 9/15
80/80 [=====] - 1s 8ms/step - loss: 6.6588e-04 - accuracy: 1.0000 - val_loss: 0.0233 - val_accuracy: 0.9949
Epoch 10/15
80/80 [=====] - 1s 8ms/step - loss: 5.5241e-04 - accuracy: 1.0000 - val_loss: 0.0226 - val_accuracy: 0.9949
Epoch 11/15
80/80 [=====] - 1s 7ms/step - loss: 4.8229e-04 - accuracy: 1.0000 - val_loss: 0.0222 - val_accuracy: 0.9949
Epoch 12/15
80/80 [=====] - 1s 8ms/step - loss: 4.2143e-04 - accuracy: 1.0000 - val_loss: 0.0217 - val_accuracy: 0.9949
Epoch 13/15
80/80 [=====] - 1s 8ms/step - loss: 3.7001e-04 - accuracy: 1.0000 - val_loss: 0.0213 - val_accuracy: 0.9949
Epoch 14/15
80/80 [=====] - 1s 7ms/step - loss: 3.2739e-04 - accuracy: 1.0000 - val_loss: 0.0214 - val_accuracy: 0.9949
Epoch 15/15
80/80 [=====] - 1s 7ms/step - loss: 2.9645e-04 - accuracy: 1.0000 - val_loss: 0.0211 - val_accuracy: 0.9949
<keras.callbacks.History at 0x7f4be9474c50>
```

```
loss: 2.9645e-04 -accuracy: 1.0 -val_loss: 0.0211 -val_accuracy: 0.9949
```

3- Batch_size=15

```
model.fit(X_train_r, y_train, epochs=15, validation_data=(X_valid_r, y_valid), batch_size=15)
```

```
Epoch 1/15
53/53 [=====] - 2s 18ms/step - loss: 1.5921 - accuracy: 0.6604 - val_loss: 0.1887 - val_accuracy: 0.9545
Epoch 2/15
53/53 [=====] - 1s 12ms/step - loss: 0.0831 - accuracy: 0.9785 - val_loss: 0.0711 - val_accuracy: 0.9949
Epoch 3/15
53/53 [=====] - 1s 15ms/step - loss: 0.0159 - accuracy: 0.9962 - val_loss: 0.0349 - val_accuracy: 1.0000
Epoch 4/15
53/53 [=====] - 1s 10ms/step - loss: 0.0042 - accuracy: 1.0000 - val_loss: 0.0375 - val_accuracy: 0.9949
Epoch 5/15
53/53 [=====] - 0s 8ms/step - loss: 0.0034 - accuracy: 0.9987 - val_loss: 0.0353 - val_accuracy: 0.9949
Epoch 6/15
53/53 [=====] - 0s 9ms/step - loss: 0.0131 - accuracy: 0.9975 - val_loss: 0.0337 - val_accuracy: 0.9899
Epoch 7/15
53/53 [=====] - 0s 9ms/step - loss: 0.0024 - accuracy: 0.9987 - val_loss: 0.0699 - val_accuracy: 0.9949
Epoch 8/15
53/53 [=====] - 0s 8ms/step - loss: 0.0017 - accuracy: 1.0000 - val_loss: 0.0478 - val_accuracy: 0.9949
Epoch 9/15
53/53 [=====] - 0s 8ms/step - loss: 8.3680e-04 - accuracy: 1.0000 - val_loss: 0.0430 - val_accuracy: 0.9949
Epoch 10/15
53/53 [=====] - 0s 8ms/step - loss: 6.6375e-04 - accuracy: 1.0000 - val_loss: 0.0407 - val_accuracy: 0.9949
Epoch 11/15
53/53 [=====] - 0s 8ms/step - loss: 5.7865e-04 - accuracy: 1.0000 - val_loss: 0.0390 - val_accuracy: 0.9949
Epoch 12/15
53/53 [=====] - 0s 8ms/step - loss: 4.9938e-04 - accuracy: 1.0000 - val_loss: 0.0379 - val_accuracy: 0.9949
Epoch 13/15
53/53 [=====] - 0s 8ms/step - loss: 4.4032e-04 - accuracy: 1.0000 - val_loss: 0.0367 - val_accuracy: 0.9949
Epoch 14/15
53/53 [=====] - 0s 8ms/step - loss: 3.9094e-04 - accuracy: 1.0000 - val_loss: 0.0361 - val_accuracy: 0.9949
Epoch 15/15
53/53 [=====] - 0s 8ms/step - loss: 3.5427e-04 - accuracy: 1.0000 - val_loss: 0.0354 - val_accuracy: 0.9949
```

```
loss: 3.5427e-04 - accuracy: 1.0000 - val_loss: 0.0354 - val_accuracy: 0.9949
```

The best batch size was 10 so we stick with it.

Next we go with hidden nodes Size. We chose the next values:

1- Hidden nodes= 512

```
model.fit(X_train_r, y_train, epochs=15, validation_data=(X_valid_r, y_valid), batch_size=15)

Epoch 1/15
53/53 [=====] - 2s 18ms/step - loss: 1.5921 - accuracy: 0.6604 - val_loss: 0.1887 - val_accuracy: 0.9545
Epoch 2/15
53/53 [=====] - 1s 12ms/step - loss: 0.0831 - accuracy: 0.9785 - val_loss: 0.0711 - val_accuracy: 0.9949
Epoch 3/15
53/53 [=====] - 1s 15ms/step - loss: 0.0159 - accuracy: 0.9962 - val_loss: 0.0349 - val_accuracy: 1.0000
Epoch 4/15
53/53 [=====] - 1s 10ms/step - loss: 0.0042 - accuracy: 1.0000 - val_loss: 0.0375 - val_accuracy: 0.9949
Epoch 5/15
53/53 [=====] - 0s 8ms/step - loss: 0.0034 - accuracy: 0.9987 - val_loss: 0.0353 - val_accuracy: 0.9949
Epoch 6/15
53/53 [=====] - 0s 9ms/step - loss: 0.0131 - accuracy: 0.9975 - val_loss: 0.0337 - val_accuracy: 0.9899
Epoch 7/15
53/53 [=====] - 0s 9ms/step - loss: 0.0024 - accuracy: 0.9987 - val_loss: 0.0699 - val_accuracy: 0.9949
Epoch 8/15
53/53 [=====] - 0s 8ms/step - loss: 0.0017 - accuracy: 1.0000 - val_loss: 0.0478 - val_accuracy: 0.9949
Epoch 9/15
53/53 [=====] - 0s 8ms/step - loss: 8.3680e-04 - accuracy: 1.0000 - val_loss: 0.0430 - val_accuracy: 0.9949
Epoch 10/15
53/53 [=====] - 0s 8ms/step - loss: 6.6375e-04 - accuracy: 1.0000 - val_loss: 0.0407 - val_accuracy: 0.9949
Epoch 11/15
53/53 [=====] - 0s 8ms/step - loss: 5.7865e-04 - accuracy: 1.0000 - val_loss: 0.0390 - val_accuracy: 0.9949
Epoch 12/15
53/53 [=====] - 0s 8ms/step - loss: 4.9938e-04 - accuracy: 1.0000 - val_loss: 0.0379 - val_accuracy: 0.9949
Epoch 13/15
53/53 [=====] - 0s 8ms/step - loss: 4.4032e-04 - accuracy: 1.0000 - val_loss: 0.0367 - val_accuracy: 0.9949
Epoch 14/15
53/53 [=====] - 0s 8ms/step - loss: 3.9094e-04 - accuracy: 1.0000 - val_loss: 0.0361 - val_accuracy: 0.9949
Epoch 15/15
53/53 [=====] - 0s 8ms/step - loss: 3.5427e-04 - accuracy: 1.0000 - val_loss: 0.0354 - val_accuracy: 0.9949
```

```
loss: 3.5427e-04 -accuracy: 1.0000 -val_loss: 0.0354 -val accuracy: 0.9949
```

2- Hidden nodes= 256

```
Epoch 1/15
53/53 [=====] - 1s 10ms/step - loss: 1.6939 - accuracy: 0.6477 - val_loss: 0.3139 - val_accuracy: 0.9293
Epoch 2/15
53/53 [=====] - 0s 6ms/step - loss: 0.0814 - accuracy: 0.9861 - val_loss: 0.1383 - val_accuracy: 0.9798
Epoch 3/15
53/53 [=====] - 0s 6ms/step - loss: 0.0224 - accuracy: 0.9962 - val_loss: 0.0565 - val_accuracy: 0.9949
Epoch 4/15
53/53 [=====] - 0s 6ms/step - loss: 0.0085 - accuracy: 0.9987 - val_loss: 0.0436 - val_accuracy: 0.9949
Epoch 5/15
53/53 [=====] - 0s 6ms/step - loss: 0.0061 - accuracy: 0.9987 - val_loss: 0.0403 - val_accuracy: 0.9949
Epoch 6/15
53/53 [=====] - 0s 7ms/step - loss: 0.0049 - accuracy: 0.9987 - val_loss: 0.0359 - val_accuracy: 0.9949
Epoch 7/15
53/53 [=====] - 0s 6ms/step - loss: 0.0022 - accuracy: 1.0000 - val_loss: 0.0344 - val_accuracy: 0.9949
Epoch 8/15
53/53 [=====] - 0s 6ms/step - loss: 0.0017 - accuracy: 1.0000 - val_loss: 0.0319 - val_accuracy: 0.9949
Epoch 9/15
53/53 [=====] - 0s 6ms/step - loss: 0.0014 - accuracy: 1.0000 - val_loss: 0.0299 - val_accuracy: 0.9949
Epoch 10/15
53/53 [=====] - 0s 6ms/step - loss: 0.0012 - accuracy: 1.0000 - val_loss: 0.0305 - val_accuracy: 0.9949
Epoch 11/15
53/53 [=====] - 0s 6ms/step - loss: 0.0010 - accuracy: 1.0000 - val_loss: 0.0290 - val_accuracy: 0.9949
Epoch 12/15
53/53 [=====] - 0s 6ms/step - loss: 8.9673e-04 - accuracy: 1.0000 - val_loss: 0.0273 - val_accuracy: 0.9949
Epoch 13/15
53/53 [=====] - 0s 6ms/step - loss: 8.0757e-04 - accuracy: 1.0000 - val_loss: 0.0271 - val_accuracy: 0.9949
Epoch 14/15
53/53 [=====] - 0s 6ms/step - loss: 7.0315e-04 - accuracy: 1.0000 - val_loss: 0.0273 - val_accuracy: 0.9949
Epoch 15/15
53/53 [=====] - 0s 7ms/step - loss: 6.3737e-04 - accuracy: 1.0000 - val_loss: 0.0273 - val_accuracy: 0.9949
<keras.callbacks.History at 0x7f4be8cfe7d0>
```

```
loss: 6.3737e-04 -accuracy: 1.00 -val_loss: 0.0273 -val_accuracy:0.9949
```

3- Hidden nodes= 1024

```
Epoch 1/15
53/53 [=====] - 2s 21ms/step - loss: 1.5200 - accuracy: 0.6730 - val_loss: 0.2454 - val_accuracy: 0.9293
Epoch 2/15
53/53 [=====] - 1s 13ms/step - loss: 0.0918 - accuracy: 0.9823 - val_loss: 0.1160 - val_accuracy: 0.9545
Epoch 3/15
53/53 [=====] - 1s 10ms/step - loss: 0.0291 - accuracy: 0.9962 - val_loss: 0.0463 - val_accuracy: 0.9899
Epoch 4/15
53/53 [=====] - 1s 10ms/step - loss: 0.0030 - accuracy: 1.0000 - val_loss: 0.0367 - val_accuracy: 0.9899
Epoch 5/15
53/53 [=====] - 1s 10ms/step - loss: 0.0015 - accuracy: 1.0000 - val_loss: 0.0275 - val_accuracy: 0.9949
Epoch 6/15
53/53 [=====] - 1s 10ms/step - loss: 8.0985e-04 - accuracy: 1.0000 - val_loss: 0.0261 - val_accuracy: 0.9949
Epoch 7/15
53/53 [=====] - 1s 10ms/step - loss: 6.4929e-04 - accuracy: 1.0000 - val_loss: 0.0245 - val_accuracy: 0.9949
Epoch 8/15
53/53 [=====] - 1s 10ms/step - loss: 5.3969e-04 - accuracy: 1.0000 - val_loss: 0.0228 - val_accuracy: 0.9949
Epoch 9/15
53/53 [=====] - 1s 10ms/step - loss: 4.4927e-04 - accuracy: 1.0000 - val_loss: 0.0210 - val_accuracy: 0.9949
Epoch 10/15
53/53 [=====] - 0s 9ms/step - loss: 3.8527e-04 - accuracy: 1.0000 - val_loss: 0.0212 - val_accuracy: 0.9949
Epoch 11/15
53/53 [=====] - 0s 9ms/step - loss: 3.3634e-04 - accuracy: 1.0000 - val_loss: 0.0209 - val_accuracy: 0.9949
Epoch 12/15
53/53 [=====] - 0s 9ms/step - loss: 2.9523e-04 - accuracy: 1.0000 - val_loss: 0.0205 - val_accuracy: 0.9949
Epoch 13/15
53/53 [=====] - 1s 10ms/step - loss: 2.6241e-04 - accuracy: 1.0000 - val_loss: 0.0201 - val_accuracy: 0.9949
Epoch 14/15
53/53 [=====] - 1s 9ms/step - loss: 2.3795e-04 - accuracy: 1.0000 - val_loss: 0.0195 - val_accuracy: 0.9949
Epoch 15/15
53/53 [=====] - 0s 9ms/step - loss: 2.1453e-04 - accuracy: 1.0000 - val_loss: 0.0198 - val_accuracy: 0.9949
<keras.callbacks.History at 0x7f4be9365450>
```

```
loss:2.1453e-04 -accuracy: 1.00 -val_loss: 0.0198 -val_accuracy: 0.9949
```

The best hidden nodes number was 1024 so we stick with it.

Next we go with dropout rate. We chose the next values:

1- Dropout= 0.2

```
Epoch 1/15
80/80 [=====] - 2s 15ms/step - loss: 1.5270 - accuracy: 0.6477 - val_loss: 0.3094 - val_accuracy: 0.9343
Epoch 2/15
80/80 [=====] - 1s 9ms/step - loss: 0.1052 - accuracy: 0.9684 - val_loss: 0.1019 - val_accuracy: 0.9697
Epoch 3/15
80/80 [=====] - 1s 10ms/step - loss: 0.0428 - accuracy: 0.9861 - val_loss: 0.1060 - val_accuracy: 0.9747
Epoch 4/15
80/80 [=====] - 1s 10ms/step - loss: 0.0491 - accuracy: 0.9899 - val_loss: 0.0935 - val_accuracy: 0.9798
Epoch 5/15
80/80 [=====] - 1s 10ms/step - loss: 0.0190 - accuracy: 0.9937 - val_loss: 0.0412 - val_accuracy: 0.9899
Epoch 6/15
80/80 [=====] - 1s 10ms/step - loss: 0.0037 - accuracy: 1.0000 - val_loss: 0.0295 - val_accuracy: 0.9949
Epoch 7/15
80/80 [=====] - 1s 9ms/step - loss: 9.3466e-04 - accuracy: 1.0000 - val_loss: 0.0256 - val_accuracy: 0.9949
Epoch 8/15
80/80 [=====] - 1s 10ms/step - loss: 3.7088e-04 - accuracy: 1.0000 - val_loss: 0.0246 - val_accuracy: 0.9949
Epoch 9/15
80/80 [=====] - 1s 9ms/step - loss: 2.9325e-04 - accuracy: 1.0000 - val_loss: 0.0236 - val_accuracy: 0.9949
Epoch 10/15
80/80 [=====] - 1s 9ms/step - loss: 2.4210e-04 - accuracy: 1.0000 - val_loss: 0.0229 - val_accuracy: 0.9949
Epoch 11/15
80/80 [=====] - 1s 9ms/step - loss: 2.0504e-04 - accuracy: 1.0000 - val_loss: 0.0224 - val_accuracy: 0.9949
Epoch 12/15
80/80 [=====] - 1s 10ms/step - loss: 1.8081e-04 - accuracy: 1.0000 - val_loss: 0.0226 - val_accuracy: 0.9949
Epoch 13/15
80/80 [=====] - 1s 10ms/step - loss: 1.6147e-04 - accuracy: 1.0000 - val_loss: 0.0213 - val_accuracy: 0.9949
Epoch 14/15
80/80 [=====] - 1s 8ms/step - loss: 1.4288e-04 - accuracy: 1.0000 - val_loss: 0.0210 - val_accuracy: 0.9949
Epoch 15/15
80/80 [=====] - 1s 9ms/step - loss: 1.2680e-04 - accuracy: 1.0000 - val_loss: 0.0208 - val_accuracy: 0.9949
<keras.callbacks.History at 0x7f4be7a5dc50>
```

```

Epoch 1/15
80/80 [=====] - 2s 19ms/step - loss: 1.5714 - accuracy: 0.6604 - val_loss: 0.3821 - val_accuracy: 0.9141
Epoch 2/15
80/80 [=====] - 1s 10ms/step - loss: 0.0903 - accuracy: 0.9798 - val_loss: 0.1124 - val_accuracy: 0.9747
Epoch 3/15
80/80 [=====] - 1s 10ms/step - loss: 0.0346 - accuracy: 0.9937 - val_loss: 0.0551 - val_accuracy: 0.9848
Epoch 4/15
80/80 [=====] - 1s 9ms/step - loss: 0.0073 - accuracy: 0.9987 - val_loss: 0.0518 - val_accuracy: 0.9848
Epoch 5/15
80/80 [=====] - 1s 10ms/step - loss: 0.0106 - accuracy: 0.9975 - val_loss: 0.0577 - val_accuracy: 0.9798
Epoch 6/15
80/80 [=====] - 1s 9ms/step - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.0479 - val_accuracy: 0.9899
Epoch 7/15
80/80 [=====] - 1s 9ms/step - loss: 5.4985e-04 - accuracy: 1.0000 - val_loss: 0.0442 - val_accuracy: 0.9949
Epoch 8/15
80/80 [=====] - 1s 9ms/step - loss: 3.8650e-04 - accuracy: 1.0000 - val_loss: 0.0420 - val_accuracy: 0.9949
Epoch 9/15
80/80 [=====] - 1s 10ms/step - loss: 3.1817e-04 - accuracy: 1.0000 - val_loss: 0.0407 - val_accuracy: 0.9949
Epoch 10/15
80/80 [=====] - 1s 9ms/step - loss: 2.6795e-04 - accuracy: 1.0000 - val_loss: 0.0393 - val_accuracy: 0.9949
Epoch 11/15
80/80 [=====] - 1s 10ms/step - loss: 2.3187e-04 - accuracy: 1.0000 - val_loss: 0.0382 - val_accuracy: 0.9949
Epoch 12/15
80/80 [=====] - 1s 10ms/step - loss: 2.0029e-04 - accuracy: 1.0000 - val_loss: 0.0372 - val_accuracy: 0.9949
Epoch 13/15
80/80 [=====] - 1s 10ms/step - loss: 1.7784e-04 - accuracy: 1.0000 - val_loss: 0.0363 - val_accuracy: 0.9949
Epoch 14/15
80/80 [=====] - 1s 9ms/step - loss: 1.5887e-04 - accuracy: 1.0000 - val_loss: 0.0356 - val_accuracy: 0.9949
Epoch 15/15
80/80 [=====] - 1s 9ms/step - loss: 1.4143e-04 - accuracy: 1.0000 - val_loss: 0.0350 - val_accuracy: 0.9949
<keras.callbacks.History at 0x7f4be916c790>

```

3- Dropout= 0.1

```

Epoch 1/15
80/80 [=====] - 1s 12ms/step - loss: 1.5127 - accuracy: 0.6730 - val_loss: 0.3080 - val_accuracy: 0.9293
Epoch 2/15
80/80 [=====] - 1s 9ms/step - loss: 0.0920 - accuracy: 0.9773 - val_loss: 0.1904 - val_accuracy: 0.9646
Epoch 3/15
80/80 [=====] - 1s 9ms/step - loss: 0.0416 - accuracy: 0.9886 - val_loss: 0.0736 - val_accuracy: 0.9899
Epoch 4/15
80/80 [=====] - 1s 10ms/step - loss: 0.0682 - accuracy: 0.9886 - val_loss: 0.0520 - val_accuracy: 0.9899
Epoch 5/15
80/80 [=====] - 1s 9ms/step - loss: 0.0552 - accuracy: 0.9937 - val_loss: 0.0763 - val_accuracy: 0.9646
Epoch 6/15
80/80 [=====] - 1s 8ms/step - loss: 0.0290 - accuracy: 0.9962 - val_loss: 0.1936 - val_accuracy: 0.9848
Epoch 7/15
80/80 [=====] - 1s 9ms/step - loss: 0.0549 - accuracy: 0.9924 - val_loss: 0.0673 - val_accuracy: 0.9747
Epoch 8/15
80/80 [=====] - 1s 8ms/step - loss: 0.1241 - accuracy: 0.9722 - val_loss: 0.2474 - val_accuracy: 0.9394
Epoch 9/15
80/80 [=====] - 1s 10ms/step - loss: 0.1692 - accuracy: 0.9545 - val_loss: 0.2165 - val_accuracy: 0.9293
Epoch 10/15
80/80 [=====] - 1s 8ms/step - loss: 0.0654 - accuracy: 0.9773 - val_loss: 0.1660 - val_accuracy: 0.9596
Epoch 11/15
80/80 [=====] - 1s 8ms/step - loss: 0.1403 - accuracy: 0.9735 - val_loss: 0.2382 - val_accuracy: 0.9293
Epoch 12/15
80/80 [=====] - 1s 10ms/step - loss: 0.0966 - accuracy: 0.9747 - val_loss: 0.3316 - val_accuracy: 0.9192
Epoch 13/15
80/80 [=====] - 1s 8ms/step - loss: 0.0544 - accuracy: 0.9886 - val_loss: 0.1426 - val_accuracy: 0.9495
Epoch 14/15
80/80 [=====] - 1s 9ms/step - loss: 0.0475 - accuracy: 0.9912 - val_loss: 0.1180 - val_accuracy: 0.9646
Epoch 15/15
80/80 [=====] - 1s 9ms/step - loss: 0.0575 - accuracy: 0.9924 - val_loss: 0.1904 - val_accuracy: 0.9646

```

```
loss: 0.0575 - accuracy: 0.9924 - val_loss: 0.1904 - val_accuracy: 0.9646
```

The best Dropout out rate was 0.2 so we stick with it.

Next we go with Optimizers. We chose the next values:

1- Adam

```

Epoch 1/15
80/80 [=====] - 1s 11ms/step - loss: 1.5383 - accuracy: 0.6768 - val_loss: 0.2485 - val_accuracy: 0.9444
Epoch 2/15
80/80 [=====] - 1s 9ms/step - loss: 0.8883 - accuracy: 0.9684 - val_loss: 0.1176 - val_accuracy: 0.9747
Epoch 3/15
80/80 [=====] - 1s 9ms/step - loss: 0.0596 - accuracy: 0.9823 - val_loss: 0.1262 - val_accuracy: 0.9646
Epoch 4/15
80/80 [=====] - 1s 9ms/step - loss: 0.0267 - accuracy: 0.9924 - val_loss: 0.1209 - val_accuracy: 0.9747
Epoch 5/15
80/80 [=====] - 1s 10ms/step - loss: 0.0065 - accuracy: 0.9987 - val_loss: 0.0812 - val_accuracy: 0.9798
Epoch 6/15
80/80 [=====] - 1s 9ms/step - loss: 0.0030 - accuracy: 0.9987 - val_loss: 0.1553 - val_accuracy: 0.9798
Epoch 7/15
80/80 [=====] - 1s 9ms/step - loss: 0.0095 - accuracy: 0.9975 - val_loss: 0.0450 - val_accuracy: 0.9848
Epoch 8/15
80/80 [=====] - 1s 10ms/step - loss: 7.2978e-04 - accuracy: 1.0000 - val_loss: 0.0310 - val_accuracy: 0.9899
Epoch 9/15
80/80 [=====] - 1s 10ms/step - loss: 3.6192e-04 - accuracy: 1.0000 - val_loss: 0.0275 - val_accuracy: 0.9949
Epoch 10/15
80/80 [=====] - 1s 9ms/step - loss: 2.9777e-04 - accuracy: 1.0000 - val_loss: 0.0249 - val_accuracy: 0.9949
Epoch 11/15
80/80 [=====] - 1s 9ms/step - loss: 2.4698e-04 - accuracy: 1.0000 - val_loss: 0.0232 - val_accuracy: 0.9949
Epoch 12/15
80/80 [=====] - 1s 9ms/step - loss: 2.0891e-04 - accuracy: 1.0000 - val_loss: 0.0215 - val_accuracy: 0.9949
Epoch 13/15
80/80 [=====] - 1s 9ms/step - loss: 1.8790e-04 - accuracy: 1.0000 - val_loss: 0.0206 - val_accuracy: 0.9949
Epoch 14/15
80/80 [=====] - 1s 10ms/step - loss: 1.6279e-04 - accuracy: 1.0000 - val_loss: 0.0196 - val_accuracy: 0.9949
Epoch 15/15
80/80 [=====] - 1s 10ms/step - loss: 1.4503e-04 - accuracy: 1.0000 - val_loss: 0.0189 - val_accuracy: 0.9949

```

```

loss: 1.4503e-04 -accuracy: 1.0000 -val_loss: 0.0189 -val_accuracy: 0.9949

```

2- SGD

```

Epoch 1/15
80/80 [=====] - 3s 19ms/step - loss: 3.7474 - accuracy: 0.3763 - val_loss: 2.7428 - val_accuracy: 0.7980
Epoch 2/15
80/80 [=====] - 1s 12ms/step - loss: 2.0615 - accuracy: 0.8447 - val_loss: 1.4185 - val_accuracy: 0.9192
Epoch 3/15
80/80 [=====] - 1s 14ms/step - loss: 0.9659 - accuracy: 0.9583 - val_loss: 0.7607 - val_accuracy: 0.9394
Epoch 4/15
80/80 [=====] - 1s 15ms/step - loss: 0.4834 - accuracy: 0.9874 - val_loss: 0.4730 - val_accuracy: 0.9848
Epoch 5/15
80/80 [=====] - 1s 10ms/step - loss: 0.2870 - accuracy: 0.9962 - val_loss: 0.3386 - val_accuracy: 0.9848
Epoch 6/15
80/80 [=====] - 1s 9ms/step - loss: 0.1924 - accuracy: 0.9975 - val_loss: 0.2665 - val_accuracy: 0.9848
Epoch 7/15
80/80 [=====] - 1s 9ms/step - loss: 0.1420 - accuracy: 0.9987 - val_loss: 0.2224 - val_accuracy: 0.9848
Epoch 8/15
80/80 [=====] - 1s 8ms/step - loss: 0.1103 - accuracy: 0.9987 - val_loss: 0.1914 - val_accuracy: 0.9848
Epoch 9/15
80/80 [=====] - 1s 8ms/step - loss: 0.0895 - accuracy: 0.9987 - val_loss: 0.1711 - val_accuracy: 0.9798
Epoch 10/15
80/80 [=====] - 1s 9ms/step - loss: 0.0761 - accuracy: 0.9987 - val_loss: 0.1562 - val_accuracy: 0.9848
Epoch 11/15
80/80 [=====] - 1s 8ms/step - loss: 0.0636 - accuracy: 1.0000 - val_loss: 0.1430 - val_accuracy: 0.9848
Epoch 12/15
80/80 [=====] - 1s 9ms/step - loss: 0.0560 - accuracy: 0.9987 - val_loss: 0.1322 - val_accuracy: 0.9798
Epoch 13/15
80/80 [=====] - 1s 8ms/step - loss: 0.0497 - accuracy: 1.0000 - val_loss: 0.1280 - val_accuracy: 0.9848
Epoch 14/15
80/80 [=====] - 1s 9ms/step - loss: 0.0451 - accuracy: 0.9987 - val_loss: 0.1172 - val_accuracy: 0.9848
Epoch 15/15
80/80 [=====] - 1s 9ms/step - loss: 0.0402 - accuracy: 1.0000 - val_loss: 0.1120 - val_accuracy: 0.9899

```

```

loss: 0.0402 - accuracy: 1.0000 - val_loss: 0.1120 - val_accuracy: 0.9899

```

3- RMSProp

```
Epoch 1/15
80/80 [=====] - 3s 20ms/step - loss: 1.6073 - accuracy: 0.6604 - val_loss: 0.2797 - val_accuracy: 0.9242
Epoch 2/15
80/80 [=====] - 1s 14ms/step - loss: 0.1214 - accuracy: 0.9672 - val_loss: 0.1973 - val_accuracy: 0.9293
Epoch 3/15
80/80 [=====] - 1s 11ms/step - loss: 0.0451 - accuracy: 0.9912 - val_loss: 0.0932 - val_accuracy: 0.9697
Epoch 4/15
80/80 [=====] - 1s 11ms/step - loss: 0.0307 - accuracy: 0.9912 - val_loss: 0.0610 - val_accuracy: 0.9798
Epoch 5/15
80/80 [=====] - 1s 12ms/step - loss: 0.0104 - accuracy: 0.9949 - val_loss: 0.1355 - val_accuracy: 0.9747
Epoch 6/15
80/80 [=====] - 1s 12ms/step - loss: 0.0016 - accuracy: 0.9987 - val_loss: 0.0588 - val_accuracy: 0.9899
Epoch 7/15
80/80 [=====] - 1s 12ms/step - loss: 0.0047 - accuracy: 0.9975 - val_loss: 0.0243 - val_accuracy: 0.9949
Epoch 8/15
80/80 [=====] - 1s 11ms/step - loss: 0.0047 - accuracy: 0.9975 - val_loss: 0.0284 - val_accuracy: 0.9848
Epoch 9/15
80/80 [=====] - 1s 12ms/step - loss: 2.6778e-05 - accuracy: 1.0000 - val_loss: 0.0561 - val_accuracy: 0.9949
Epoch 10/15
80/80 [=====] - 1s 12ms/step - loss: 3.3315e-06 - accuracy: 1.0000 - val_loss: 0.0051 - val_accuracy: 1.0000
Epoch 11/15
80/80 [=====] - 1s 12ms/step - loss: 3.3627e-06 - accuracy: 1.0000 - val_loss: 0.1094 - val_accuracy: 0.9899
Epoch 12/15
80/80 [=====] - 1s 12ms/step - loss: 4.7571e-06 - accuracy: 1.0000 - val_loss: 0.0191 - val_accuracy: 0.9899
Epoch 13/15
80/80 [=====] - 1s 12ms/step - loss: 9.3771e-08 - accuracy: 1.0000 - val_loss: 0.0893 - val_accuracy: 0.9949
Epoch 14/15
80/80 [=====] - 1s 11ms/step - loss: 5.0819e-06 - accuracy: 1.0000 - val_loss: 0.0722 - val_accuracy: 0.9899
Epoch 15/15
80/80 [=====] - 1s 11ms/step - loss: 0.0032 - accuracy: 0.9987 - val_loss: 0.0268 - val_accuracy: 0.9899
```

```
loss: 0.0032 -accuracy: 0.9987 -val_loss: 0.0268 -val_accuracy: 0.9899
```

The best Optimizer was Adam so we stick with it.